

# Delinquency Telecom Model

## Definition:

- **Delinquency** is a condition that arises when an activity or situation does not occur at its scheduled (or expected) date i.e., it occurs later than expected.

## Use Case:

- Many donors, experts, and microfinance institutions (MFI) have become convinced that using mobile financial services (MFS) is more convenient and efficient, and less costly, than the traditional high-touch model for delivering microfinance services. MFS becomes especially useful when targeting the unbanked poor living in remote areas. The implementation of MFS, though, has been uneven with both significant challenges and successes.
- Today, microfinance is widely accepted as a poverty-reduction tool, representing \$70 billion in outstanding loans and a global outreach of 200 million clients.
- One of our Client in Telecom collaborates with an MFI to provide micro-credit on mobile balances to be paid back in 5 days. The Consumer is believed to be delinquent if he deviates from the path of paying back the loaned amount within 5 days

## Machine Learning problem :

- Create a delinquency model which can predict in terms of a probability for each loan transaction, whether the customer will be paying back the loaned amount within 5 days of insurance of loan (Label '1' & '0')
- Basically a Binary Classification setup

## Real-world/Business objectives and constraints.

- No low-latency requirement.
- Interpretability is important.
- Probability of a data-point belonging to each class is needed.

## Performance Metric

- Log-loss (Since probabilities is our concern)
- Confusion matrix (Also want to check some precision and recalls)

In [1]:

```
cd drive/My Drive/Algo8
```

```
/content/drive/My Drive/Algo8
```

In [3]:

```
!pip install catboost
```

Collecting catboost

Downloading

[https://files.pythonhosted.org/packages/b2/aa/e61819d04ef2bbee778bf4b3a748db1f3ad23512377e43ecfdc327a0/catboost-0.23.2-cp36-none-manylinux1\\_x86\\_64.whl](https://files.pythonhosted.org/packages/b2/aa/e61819d04ef2bbee778bf4b3a748db1f3ad23512377e43ecfdc327a0/catboost-0.23.2-cp36-none-manylinux1_x86_64.whl) (64.8MB)

|██| 64.8MB 49kB/s

Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (from catboost) (1.4.1)

Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from catboost) (1.12.0)

Requirement already satisfied: graphviz in /usr/local/lib/python3.6/dist-packages (from catboost) (0.10.1)

Requirement already satisfied: pandas>=0.24.0 in /usr/local/lib/python3.6/dist-packages (from catboost) (1.0.5)

Requirement already satisfied: plotly in /usr/local/lib/python3.6/dist-packages (from catboost) (4.4.1)

Requirement already satisfied: numpy>=1.16.0 in /usr/local/lib/python3.6/dist-packages (from catboost) (1.18.5)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.6/dist-packages (from catboost) (3.2.2)

Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.6/dist-packages (from pandas>=0.24.0->catboost) (2018.9)  
Requirement already satisfied: python-dateutil>=2.6.1 in /usr/local/lib/python3.6/dist-packages (from pandas>=0.24.0->catboost) (2.8.1)  
Requirement already satisfied: retrying>=1.3.3 in /usr/local/lib/python3.6/dist-packages (from plotly->catboost) (1.3.3)  
Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.6/dist-packages (from matplotlib->catboost) (0.10.0)  
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->catboost) (2.4.7)  
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->catboost) (1.2.0)  
Installing collected packages: catboost  
Successfully installed catboost-0.23.2

In [4]:

```
import numpy as np
import pandas as pd
import random
import seaborn as sns
import matplotlib.pyplot as plt
import pickle
%matplotlib inline
sns.set(color_codes=True)
import os
from sklearn.model_selection import GridSearchCV
from datetime import datetime
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix

from sklearn.preprocessing import StandardScaler

from sklearn import tree
from sklearn import metrics
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier

# Boosting Algorithms :
from xgboost import XGBClassifier
from catboost import CatBoostClassifier
from lightgbm import LGBMClassifier
from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.calibration import CalibratedClassifierCV
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split, GridSearchCV, StratifiedKFold
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.multiclass import OneVsRestClassifier
from sklearn.metrics import confusion_matrix, normalized_mutual_info_score
from sklearn.linear_model import SGDClassifier
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/utils/deprecation.py:144: FutureWarning: The
sklearn.metrics.classification module is deprecated in version 0.22 and will be removed in
version 0.24. The corresponding classes / functions should instead be imported from
sklearn.metrics. Anything that cannot be imported from sklearn.metrics is now part of the private
API.
  warnings.warn(message, FutureWarning)
```

In [5]:

```
train = pd.read_csv('sample_data_intw.csv')
```

## Exploratory Data Analysis and Data Preprocessing

In [6]:

```
train.head()
```

Out[6]:

Unnamed: 0	label	msisdn	aon	daily_decr30	daily_decr90	rental30	rental90	last_rech_date_ma	last_rech_date_da	last_rec
0	1	0	21408I70789	272.0	3055.050000	3065.150000	220.13	260.13	2.0	0.0
1	2	1	76462I70374	712.0	12122.000000	12124.750000	3691.26	3691.26	20.0	0.0
2	3	1	17943I70372	535.0	1398.000000	1398.000000	900.13	900.13	3.0	0.0
3	4	1	55773I70781	241.0	21.228000	21.228000	159.42	159.42	41.0	0.0
4	5	1	03813I82730	947.0	150.619333	150.619333	1098.90	1098.90	4.0	0.0

In [7]:

```
train.drop('Unnamed: 0',axis=1,inplace=True)
```

In [8]:

```
print("Size of Train data = {}".format(train.shape))
```

Size of Train data = (209593, 36)

## Checks for Null values

In [9]:

```
train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209593 entries, 0 to 209592
Data columns (total 36 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   label                                209593 non-null  int64
 1   msisdn                               209593 non-null  object
 2   aon                                   209593 non-null  float64
 3   daily_decr30                         209593 non-null  float64
 4   daily_decr90                         209593 non-null  float64
 5   rental30                             209593 non-null  float64
 6   rental90                             209593 non-null  float64
 7   last_rech_date_ma                    209593 non-null  float64
 8   last_rech_date_da                    209593 non-null  float64
 9   last_rech_amt_ma                     209593 non-null  int64
10  cnt_ma_rech30                        209593 non-null  int64
11  fr_ma_rech30                         209593 non-null  float64
12  sumamnt_ma_rech30                    209593 non-null  float64
13  medianamnt_ma_rech30                 209593 non-null  float64
14  medianmarechprebal30                 209593 non-null  float64
15  cnt_ma_rech90                        209593 non-null  int64
16  fr_ma_rech90                         209593 non-null  int64
17  sumamnt_ma_rech90                    209593 non-null  int64
18  medianamnt_ma_rech90                 209593 non-null  float64
19  medianmarechprebal90                 209593 non-null  float64
20  cnt_da_rech30                        209593 non-null  float64
21  fr_da_rech30                         209593 non-null  float64
22  cnt_da_rech90                        209593 non-null  int64
23  fr_da_rech90                         209593 non-null  int64
24  cnt_loans30                          209593 non-null  int64
25  amnt_loans30                         209593 non-null  int64
26  maxamnt_loans30                      209593 non-null  float64
27  medianamnt_loans30                   209593 non-null  float64
28  cnt_loans90                          209593 non-null  float64
29  amnt_loans90                         209593 non-null  int64
30  maxamnt_loans90                      209593 non-null  int64
31  medianamnt_loans90                   209593 non-null  float64
32  payback30                            209593 non-null  float64
33  payback90                            209593 non-null  float64
34  pcircle                              209593 non-null  object
```

```
34  pdate          209593 non-null object
35  pdate          209593 non-null object
dtypes: float64(21), int64(12), object(3)
memory usage: 57.6+ MB
```

In [10]:

```
train.isnull().sum()
```

Out[10]:

```
label          0
msisdn         0
aon            0
daily_decr30   0
daily_decr90   0
rental30       0
rental90       0
last_rech_date_ma  0
last_rech_date_da  0
last_rech_amt_ma  0
cnt_ma_rech30   0
fr_ma_rech30   0
sumamnt_ma_rech30  0
medianamnt_ma_rech30  0
medianmarechprebal30  0
cnt_ma_rech90   0
fr_ma_rech90   0
sumamnt_ma_rech90  0
medianamnt_ma_rech90  0
medianmarechprebal90  0
cnt_da_rech30   0
fr_da_rech30   0
cnt_da_rech90   0
fr_da_rech90   0
cnt_loans30     0
amnt_loans30    0
maxamnt_loans30  0
medianamnt_loans30  0
cnt_loans90     0
amnt_loans90    0
maxamnt_loans90  0
medianamnt_loans90  0
payback30       0
payback90       0
pcircle         0
pdate          0
dtype: int64
```

In [11]:

```
train['pcircle'].value_counts()
```

Out[11]:

```
UPW    209593
Name: pcircle, dtype: int64
```

In [12]:

```
train.drop('pcircle',axis=1,inplace=True) #Same value , so not much informative
```

In [13]:

```
# Checking for duplicate values
print("Number of duplicate values in train data is "+str(sum(train.duplicated())))
```

Number of duplicate values in train data is 1

## Separating features and class labels

In [37]:

```
X = train
X = X.drop(["label"], axis = 1)

y = train['label']
```

In [117]:

```
X.shape , y.shape
```

Out[117]:


```
((209593, 32), (209593,))
```

## Checking Data Imbalances

In [39]:

```
print(train['label'].value_counts())
f,ax=plt.subplots(1,2,figsize=(16,6))
labels = ['0', '1']
train['label'].value_counts().plot.pie(explode=[0,0.2], autopct='%1.1f%%', ax=ax[0], shadow=True, label
s=labels, fontsize=10)
sns.countplot('label', data=train, ax=ax[1])
ax[1].set_xticklabels(['0', '1'], fontsize=10)
plt.show()
```

```
1    183431
0     26162
Name: label, dtype: int64
```



- Imbalanced Data

In [17]:

```
## SEE the number of outliers

Q1 = train.quantile(0.25)
Q3 = train.quantile(0.75)
IQR = Q3 - Q1
print('No. of outliers in all the fields: ', ((train < (Q1 - 1.5 * IQR)) | (train > (Q3 + 1.5 * IQR)
)).sum())
```

```
No. of outliers in all the fields:  amnt_loans30    10416
amnt_loans90    12590
age            2607
```

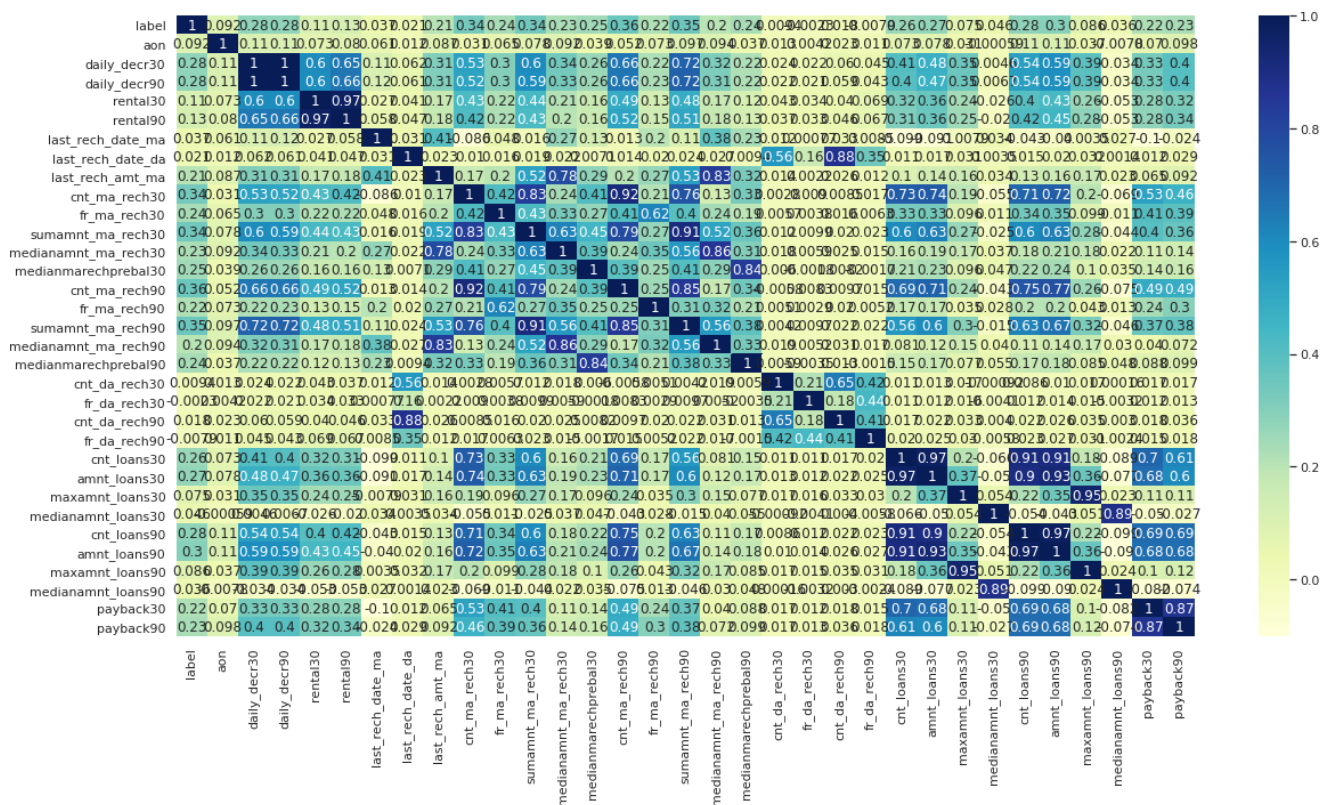
```
don 300 /
cnt_da_rech30 4114
cnt_da_rech90 5367
cnt_loans30 7817
cnt_loans90 11523
cnt_ma_rech30 11294
cnt_ma_rech90 14155
daily_decr30 16350
daily_decr90 18187
fr_da_rech30 1579
fr_da_rech90 865
fr_ma_rech30 11450
fr_ma_rech90 26845
label 26162
last_rech_amt_ma 20864
last_rech_date_da 6732
last_rech_date_ma 20145
maxamnt_loans30 30400
maxamnt_loans90 28648
medianamnt_loans30 14148
medianamnt_loans90 12169
medianamnt_ma_rech30 24928
medianamnt_ma_rech90 25457
medianmarechprebal30 27252
medianmarechprebal90 25933
msisdn 0
payback30 16532
payback90 17850
pdate 0
rental30 18526
rental90 19399
sumamnt_ma_rech30 13219
sumamnt_ma_rech90 13954
dtype: int64
```

In [18]:

```
# Correlations
f, ax = plt.subplots(figsize=(20, 10))
sns.heatmap(train.corr(method='spearman'), annot=True, cmap="YlGnBu")
```

Out[18]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f2067fba3c8>



## Convert all columns to numeric

In [40]:

```
for i in X.columns:
    if i=='pdate':
        continue
    else:
        X[i]=pd.to_numeric(X[i],errors='coerce')
```

In [41]:

```
train['msisdn'].value_counts()
```

Out[41]:

```
04581I85330    7
47819I90840    7
29191I82738    6
43430I70786    6
71742I90843    6
..
06791I70785    1
09434I82730    1
65674I70370    1
76802I89231    1
18134I85330    1
Name: msisdn, Length: 186243, dtype: int64
```

In [42]:

```
X.drop(['msisdn','pdate'],axis=1,inplace=True) # Not much informative in this case
```

In [22]:

```
X = np.array(X)
```

## Train Test Split

In [43]:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
X_train,X_cv,y_train,y_cv = train_test_split(X_train,y_train,test_size = 0.25,random_state = 42)
```

## Standardize the features

In [95]:

```
#Use standardscaler to standardize the features

sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_cv = sc.transform(X_cv)
X_test = sc.transform(X_test)
```

In [96]:

```
(len(X_train),len(y_train),len(X_test),len(y_test),len(X_cv),len(y_cv))
```

Out[96]:

```
(117895, 117895, 52399, 52399, 39299, 39299)
```

## UTILITY FUNCTIONS

In [25]:

```
def plot_matrix(matrix, labels):
    plt.figure(figsize=(20,7))
    sns.heatmap(matrix, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.show()

# This function plots the confusion matrices given y_i, y_i_hat.
def plot_confusion_matrix(test_y, predict_y):
    cm = confusion_matrix(test_y, predict_y)
    # C = 9,9 matrix, each cell (i,j) represents number of points of class i are predicted class j

    recall_table = (((cm.T)/(cm.sum(axis=1))).T)
    # How did we calculate recall_table :
    # divide each element of the confusion matrix with the sum of elements in that column
    # C = [[1, 2],
    #      [3, 4]]
    # C.T = [[1, 3],
    #        [2, 4]]
    # C.sum(axis = 1) axis=0 corresponds to columns and axis=1 corresponds to rows in two
    dimensional array
    # C.sum(axis = 1) = [[3, 7]]
    # ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7]
    #                             [2/3, 4/7]]
    # ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3]
    #                               [3/7, 4/7]]
    # sum of row elements = 1

    precision_table = (cm/cm.sum(axis=0))
    # How did we calculate precision_table :
    # divide each element of the confusion matrix with the sum of elements in that row
    # C = [[1, 2],
    #      [3, 4]]
    # C.sum(axis = 0) axis=0 corresponds to columns and axis=1 corresponds to rows in two
    dimensional array
    # C.sum(axis = 0) = [[4, 6]]
    # (C/C.sum(axis=0)) = [[1/4, 2/6],
    #                       [3/4, 4/6]]

    labels = [0,1]
    print()
    print("-"*20, "Confusion matrix", "-"*20)
    plot_matrix(cm, labels)

    print("-"*20, "Precision matrix (Column Sum=1)", "-"*20)
    plot_matrix(precision_table, labels)

    print("-"*20, "Recall matrix (Row sum=1)", "-"*20)
    plot_matrix(recall_table, labels)
```

In [32]:

```
#Data preparation for ML models.

#Misc. functions for ML models

def predict_and_plot_confusion_matrix(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    pred_y = sig_clf.predict(test_x)

    # for calculating log_loss we will provide the array of probabilities belongs to each class
    print("Log loss :", log_loss(test_y, sig_clf.predict_proba(test_x)))
    # calculating the number of data points that are misclassified
    print("Number of mis-classified points :", np.count_nonzero((pred_y- test_y))/test_y.shape[0])
    plot_confusion_matrix(test_y, pred_y)
```



In [33]:

```
def report_log_loss(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    sig_clf_probs = sig_clf.predict_proba(test_x)
    return log_loss(test_y, sig_clf_probs, eps=1e-15)
```

In [100]:

```
Xr = np.array(X_test)
yr = np.array(y_test)
```

## NOTE :

- Since we want a probabilistic interpretation from the model so we will use **LogLoss** as the Metric here

## Prediction using a 'Random' Model

- We build a random model to compare the log- loss of random model with the ML models used by us.
- In a 'Random' Model, we generate the '2' class probabilities randomly such that they sum to 1.

In [97]:

```
# We need to generate 9 numbers and the sum of numbers should be 1
# one solution is to generate 9 numbers and divide each of the numbers by their sum
# ref: https://stackoverflow.com/a/18662466/4084039
test_data_len = X_test.shape[0]
cv_data_len = X_cv.shape[0]

# we create a output array that has exactly same size as the CV data
cv_predicted_y = np.zeros((cv_data_len,2))
for i in range(cv_data_len):
    rand_probs = np.random.rand(1,2)
    cv_predicted_y[i] = ((rand_probs/sum(sum(rand_probs)))[0])
print("Log loss on Cross Validation Data using Random Model",log_loss(y_cv,cv_predicted_y, eps=1e-15))

# Test-Set error.
# We create a output array that has exactly same as the test data
test_predicted_y = np.zeros((test_data_len,2))
for i in range(test_data_len):
    rand_probs = np.random.rand(1,2)
    test_predicted_y[i] = ((rand_probs/sum(sum(rand_probs)))[0])
print("Log loss on Test Data using Random Model",log_loss(y_test,test_predicted_y, eps=1e-15))


predicted_y =np.argmax(test_predicted_y, axis=1)
plot_confusion_matrix(y_test, predicted_y)
```

Log loss on Cross Validation Data using Random Model 0.8894681457527356


Log loss on Test Data using Random Model 0.8821834697913175

----- Confusion matrix -----

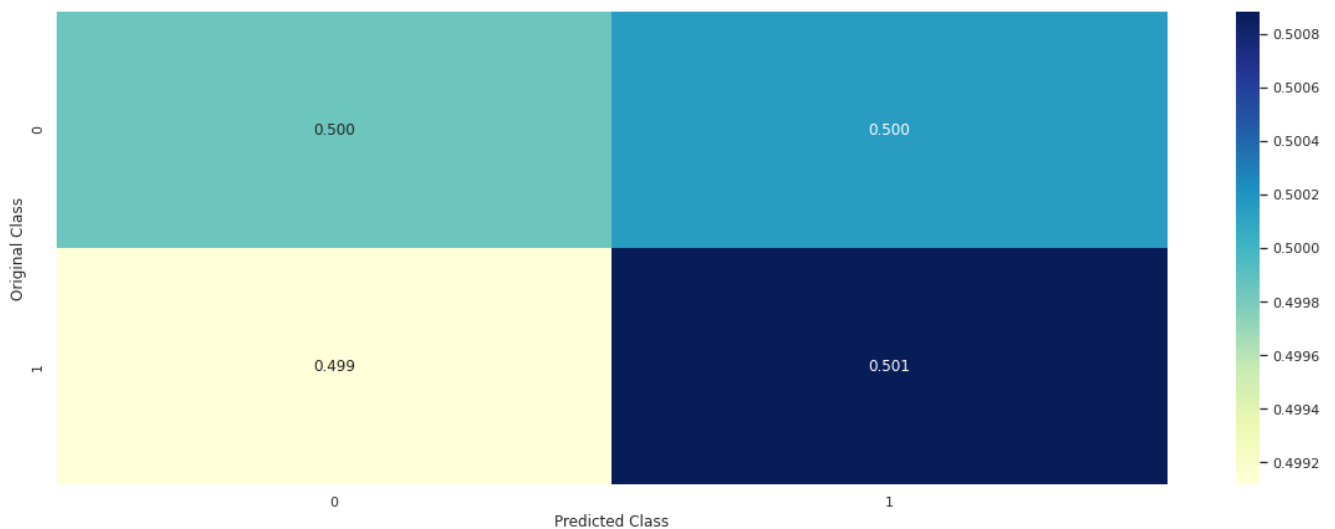




----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



## Logistic Regression with class balancing

In [99]:

```
alpha = [10 ** x for x in range(-6, 6)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(X_train,y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train,y_train)
    sig_clf_probs = sig_clf.predict_proba(X_cv)
    cv_log_error_array.append(log_loss(y_cv, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability estimates
    print("Log Loss :",log_loss(y_cv, sig_clf_probs))

fig, ax = plt.subplots()
```

```

ax.plot(alpha, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],str(txt)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(X_train,y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train,y_train)

predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ',
      alpha[best_alpha],
      "The train log loss is:",
      log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(X_cv)
print('For values of best alpha = ',
      alpha[best_alpha],
      "The cross validation log loss is:",
      log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))


predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ',
      alpha[best_alpha],
      "The test log loss is:",
      log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```

for alpha = 1e-06
Log Loss : 0.351791337006705
for alpha = 1e-05
Log Loss : 0.3320307346193461
for alpha = 0.0001
Log Loss : 0.3058021701302629
for alpha = 0.001
Log Loss : 0.2979462032962917
for alpha = 0.01
Log Loss : 0.29844309560647647
for alpha = 0.1
Log Loss : 0.30031521568767255
for alpha = 1
Log Loss : 0.30476584119358274
for alpha = 10
Log Loss : 0.30754576504364484
for alpha = 100
Log Loss : 0.3081048242033373
for alpha = 1000
Log Loss : 0.30813211508897875
for alpha = 10000
Log Loss : 0.30811283799886485
for alpha = 100000
Log Loss : 0.30870186930726273

```



For values of best alpha = 0.001 The train log loss is: 0.29879942213560184  
 For values of best alpha = 0.001 The cross validation log loss is: 0.2979462032962917  
 For values of best alpha = 0.001 The test log loss is: 0.3025684286785776

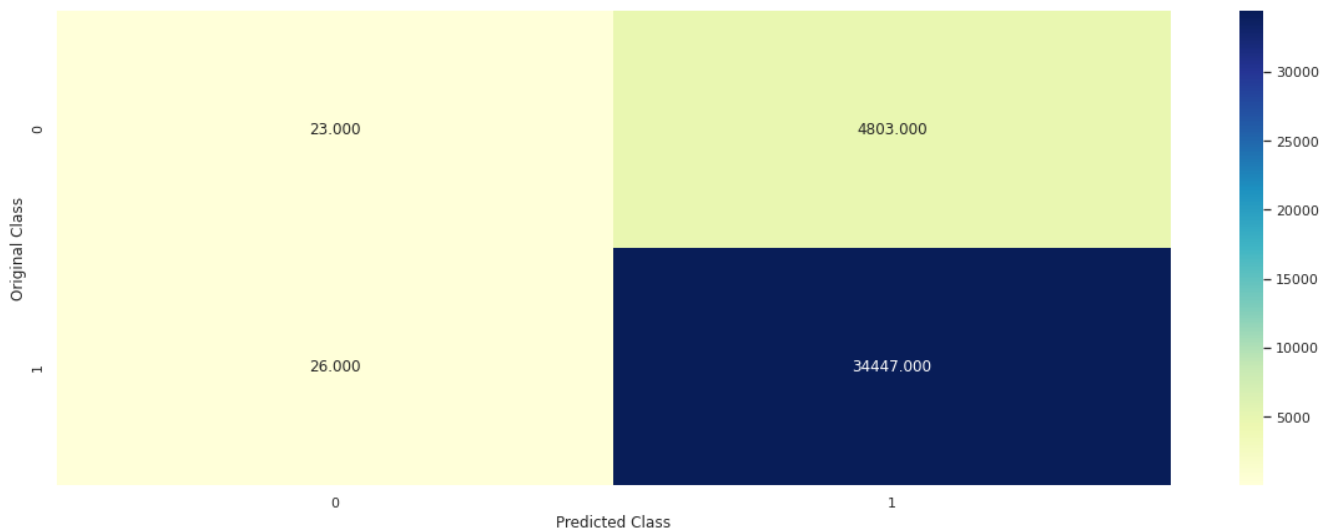
In [101]:

```
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
predict_and_plot_confusion_matrix(X_train, y_train, X_cv, y_cv, clf)
```

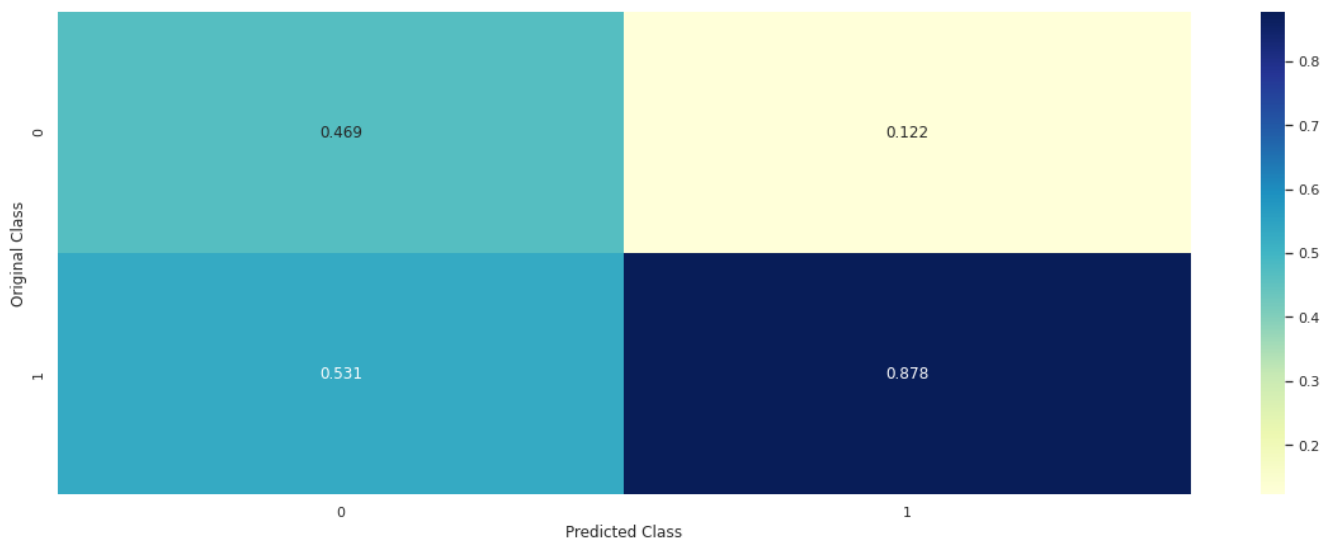
Log loss : 0.2979462032962917

Number of mis-classified points : 0.12287844474414107

----- Confusion matrix -----




----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----





## Test some points out

- Correctly predicted

In [102]:

```
# from tabulate import tabulate
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(X_train, y_train)
test_point_index = 1
no_feature = 1000
predicted_cls = sig_clf.predict(Xr[test_point_index].reshape(1, -1))
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
      np.round(sig_clf.predict_proba(Xr[test_point_index].reshape(1, -1)), 4))
print("Actual Class :", yr[test_point_index].reshape(1, -1))
indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
```

```
Predicted Class : 1
Predicted Class Probabilities: [[0.3022 0.6978]]
Actual Class : [[1]]
```

- Incorrectly predicted

In [103]:

```
# from tabulate import tabulate
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(X_train, y_train)
test_point_index = 5456
no_feature = 1000
predicted_cls = sig_clf.predict(Xr[test_point_index].reshape(1, -1))
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
      np.round(sig_clf.predict_proba(Xr[test_point_index].reshape(1, -1)), 4))
print("Actual Class :", yr[test_point_index].reshape(1, -1))
indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
```

```
Predicted Class : 1
Predicted Class Probabilities: [[0.3461 0.6539]]
Actual Class : [[0]]
```

## Linear Support Vector Machines

In [106]:

```
alpha = [10 ** x for x in range(-6, 6)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(alpha=i, penalty='l2', loss='hinge', random_state=42)
    clf.fit(X_train, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
```

```

sig_clf.fit(X_train,y_train)
sig_clf_probs = sig_clf.predict_proba(X_cv)
cv_log_error_array.append(log_loss(y_cv, sig_clf_probs, labels=clf.classes_, eps=1e-15))
# to avoid rounding error while multiplying probabilities we use log-probability estimates
print("Log Loss :",log_loss(y_cv, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],str(txt)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='hinge', random_state=42)
clf.fit(X_train,y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train,y_train)

predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ',
      alpha[best_alpha],
      "The train log loss is:",
      log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(X_cv)
print('For values of best alpha = ',
      alpha[best_alpha],
      "The cross validation log loss is:",
      log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))


predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ',
      alpha[best_alpha],
      "The test log loss is:",
      log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))


```

```

for alpha = 1e-06
Log Loss : 0.3437552901292161
for alpha = 1e-05
Log Loss : 0.3392162294716688
for alpha = 0.0001
Log Loss : 0.3372957355249959
for alpha = 0.001
Log Loss : 0.33654962358559115
for alpha = 0.01
Log Loss : 0.3406318923936145
for alpha = 0.1
Log Loss : 0.33714109225108807
for alpha = 1
Log Loss : 0.3193831797221952
for alpha = 10
Log Loss : 0.34723115385559966
for alpha = 100
Log Loss : 0.3081344831530419
for alpha = 1000
Log Loss : 0.3081429028124793
for alpha = 10000
Log Loss : 0.30809919918544515
for alpha = 100000
Log Loss : 0.30861717724991755

```





For values of best alpha = 10000 The train log loss is: 0.30915405397850865  
 For values of best alpha = 10000 The cross validation log loss is: 0.30809919918544515  
 For values of best alpha = 10000 The test log loss is: 0.3126569914606996


In [107]:

```
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='hinge', random_state=42)
predict_and_plot_confusion_matrix(X_train, y_train, X_cv, y_cv, clf)
```


Log loss : 0.30809919918544515  
 Number of mis-classified points : 0.1228021069238403

----- Confusion matrix -----

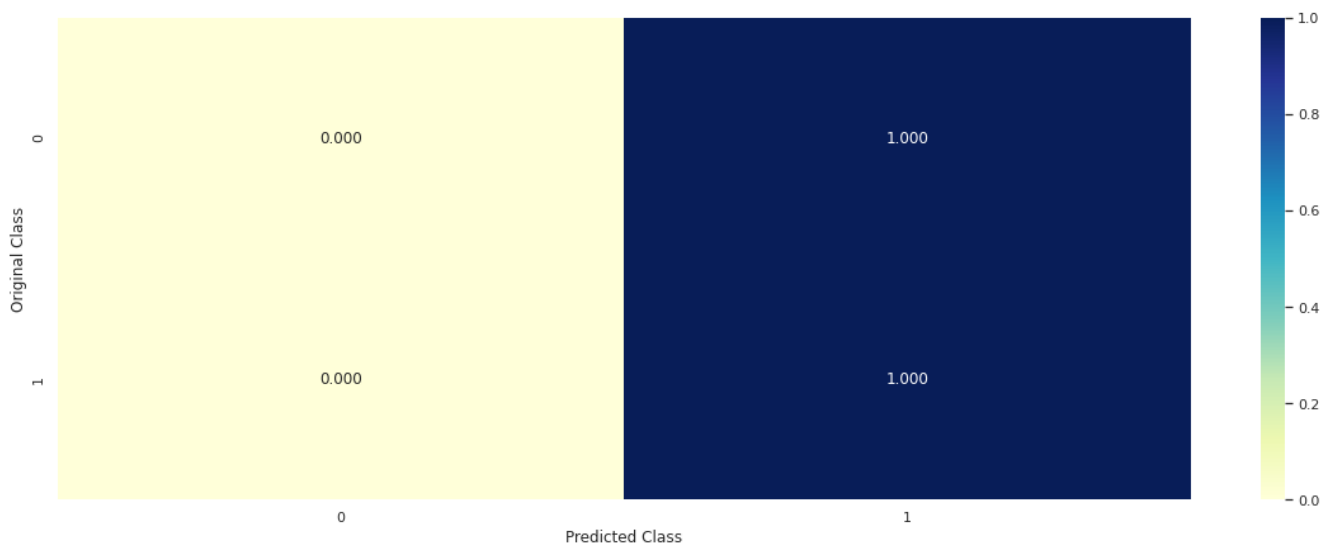
/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:28: RuntimeWarning: invalid value encountered in true\_divide



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



## Test some points out

- Correctly Classified

In [108]:

```
# from tabulate import tabulate
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='hinge', random_state=42)
clf.fit(X_train,y_train)
test_point_index = 1
no_feature = 1000
predicted_cls = sig_clf.predict(Xr[test_point_index].reshape(1, -1))
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(Xr[test_point_index].reshape(1, -1)),4))
print("Actual Class :", yr[test_point_index].reshape(1, -1))
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,no_feature]
```

Predicted Class : 1  
Predicted Class Probabilities: [[0.0735 0.9265]]  
Actual Class : [[1]]

- Incorrectly Classified

In [109]:

```
# from tabulate import tabulate
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='hinge', r
andom_state=42)
clf.fit(X_train,y_train)
test_point_index = 5456
no_feature = 1000
predicted_cls = sig_clf.predict(Xr[test_point_index].reshape(1, -1))
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(Xr[test_point_index].reshape(1, -1)),4))
print("Actual Class :", yr[test_point_index].reshape(1, -1))
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,no_feature]
```

Predicted Class : 1  
Predicted Class Probabilities: [[0.3741 0.6259]]  
Actual Class : [[0]]

## Random Forest

In [110]:



```

alpha = [100,300,500]
max_depth = [3, 5]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators =", i,"and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini', max_depth=j, random_state=42
, n_jobs=-1)
        clf.fit(X_train, y_train)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(X_train, y_train)
        sig_clf.probs = sig_clf.predict_proba(X_cv)
        cv_log_error_array.append(log_loss(y_cv, sig_clf.probs, labels=clf.classes_, eps=1e-15))
        print("Log Loss :",log_loss(y_cv, sig_clf.probs))

best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion='gini', max_depth=max
_depth[int(best_alpha*2)], random_state=42, n_jobs=-1)
clf.fit(X_train, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train, y_train)

predict_y = sig_clf.predict_proba(X_train)
print('For values of best estimator = ',
      alpha[int(best_alpha/2)],
      "The train log loss is:",
      log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(X_cv)
print('For values of best estimator = ',
      alpha[int(best_alpha/2)],
      "The cross validation log loss is:",
      log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(X_test)
print('For values of best estimator = ',
      alpha[int(best_alpha/2)],
      "The test log loss is:",
      log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```

for n_estimators = 100 and max depth = 3
Log Loss : 0.2784886966047876
for n_estimators = 100 and max depth = 5
Log Loss : 0.2669527019385133
for n_estimators = 300 and max depth = 3
Log Loss : 0.2782845046015356
for n_estimators = 300 and max depth = 5
Log Loss : 0.26715790109268056
for n_estimators = 500 and max depth = 3
Log Loss : 0.2785084940520751
for n_estimators = 500 and max depth = 5
Log Loss : 0.2672758826707215
For values of best estimator = 100 The train log loss is: 0.26520610571689845
For values of best estimator = 100 The cross validation log loss is: 0.2669527019385133
For values of best estimator = 100 The test log loss is: 0.27225935017934794

```

In [111]:

```

clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion='gini', max_depth=max
_depth[int(best_alpha*2)], random_state=42, n_jobs=-1)
predict_and_plot_confusion_matrix(X_train, y_train,X_cv,y_cv, clf)


```

```

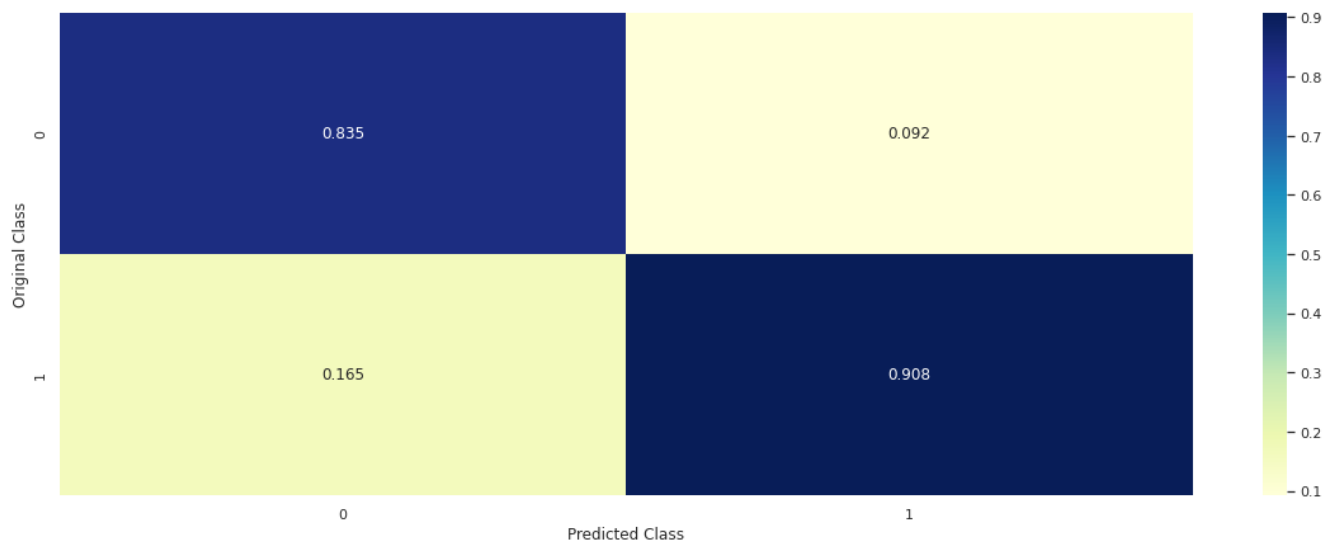
Log loss : 0.2669527019385133
Number of mis-classified points : 0.0952187078551617

```

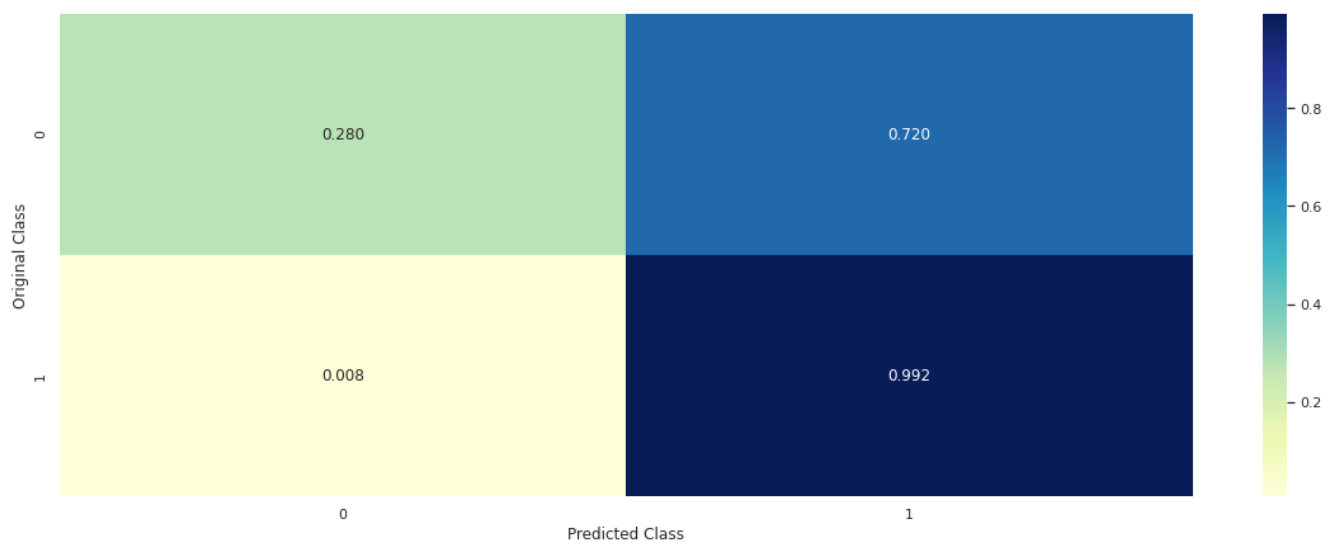
----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



## Test some points out

- Correctly classified

In [112]:

```
test_point_index = 5
no_feature = 1000
```

```

predicted_cls = sig_clf.predict(Xr[test_point_index].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(Xr[test_point_index].reshape(1,-1)),4))
print("Actual Class :", yr[test_point_index].reshape(1,-1))
indices = np.argsort(-clf.feature_importances_)

```

```

Predicted Class : 1
Predicted Class Probabilities: [[0.0428 0.9572]]
Actual Class : [[1]]

```

- Incorrectly Classified

In [113]:

```

test_point_index = 5456
no_feature = 1000
predicted_cls = sig_clf.predict(Xr[test_point_index].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(Xr[test_point_index].reshape(1,-1)),4))
print("Actual Class :", yr[test_point_index].reshape(1,-1))
indices = np.argsort(-clf.feature_importances_)

```

```

Predicted Class : 1
Predicted Class Probabilities: [[0.2557 0.7443]]
Actual Class : [[0]]

```

## Let's try UPSAMPLING

In [119]:

```

# define oversampling strategy
from imblearn.over_sampling import RandomOverSampler

oversample = RandomOverSampler(sampling_strategy='minority')
# fit and apply the transform
X_over, y_over = oversample.fit_resample(X, y)
print('Before Upsampling',X.shape, ' ', y.shape)
print('After Upsampling',X_over.shape, ' ', y_over.shape)

```

```

Before Upsampling (209593, 32)    (209593,)
After Upsampling (366862, 32)    (366862,)

```

In [120]:

```

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_over, y_over, test_size=0.25, random_state=42)
X_train,X_cv,y_train,y_cv = train_test_split(X_train,y_train,test_size = 0.25,random_state = 42)

```

In [121]:

```

#Use StandardScaler to standardize the features

sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_cv = sc.transform(X_cv)
X_test = sc.transform(X_test)

```

## Logistic Regression

In [122]:

```

alpha = [10 ** x for x in range(-6, 6)]
cv_log_error_array = []

```

```

for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log', random_state=42)

    clf.fit(X_train,y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train,y_train)
    sig_clf_probs = sig_clf.predict_proba(X_cv)
    cv_log_error_array.append(log_loss(y_cv, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability estimates
    print("Log Loss :",log_loss(y_cv, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[i],str(txt)), (alpha[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(X_train,y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train,y_train)

predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ',
      alpha[best_alpha],
      "The train log loss is:",
      log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))

predict_y = sig_clf.predict_proba(X_cv)
print('For values of best alpha = ',
      alpha[best_alpha],
      "The cross validation log loss is:",
      log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))


predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ',
      alpha[best_alpha],
      "The test log loss is:",
      log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```

for alpha = 1e-06
Log Loss : 0.618784864786861
for alpha = 1e-05
Log Loss : 0.5733931929069112
for alpha = 0.0001
Log Loss : 0.5390495776409138
for alpha = 0.001
Log Loss : 0.5217361177668973
for alpha = 0.01
Log Loss : 0.5225658118147732
for alpha = 0.1
Log Loss : 0.5282197968987324
for alpha = 1
Log Loss : 0.5377992266506653
for alpha = 10
Log Loss : 0.5435595725398746
for alpha = 100
Log Loss : 0.5445692726215877
for alpha = 1000
Log Loss : 0.544598017563088
for alpha = 10000
Log Loss : 0.5446731351013397
for alpha = 100000
Log Loss : 0.5446497846110937

```




For values of best alpha = 0.001 The train log loss is: 0.5226767926664718  
 For values of best alpha = 0.001 The cross validation log loss is: 0.5217361177668973  
 For values of best alpha = 0.001 The test log loss is: 0.5228571389582798

In [123]:

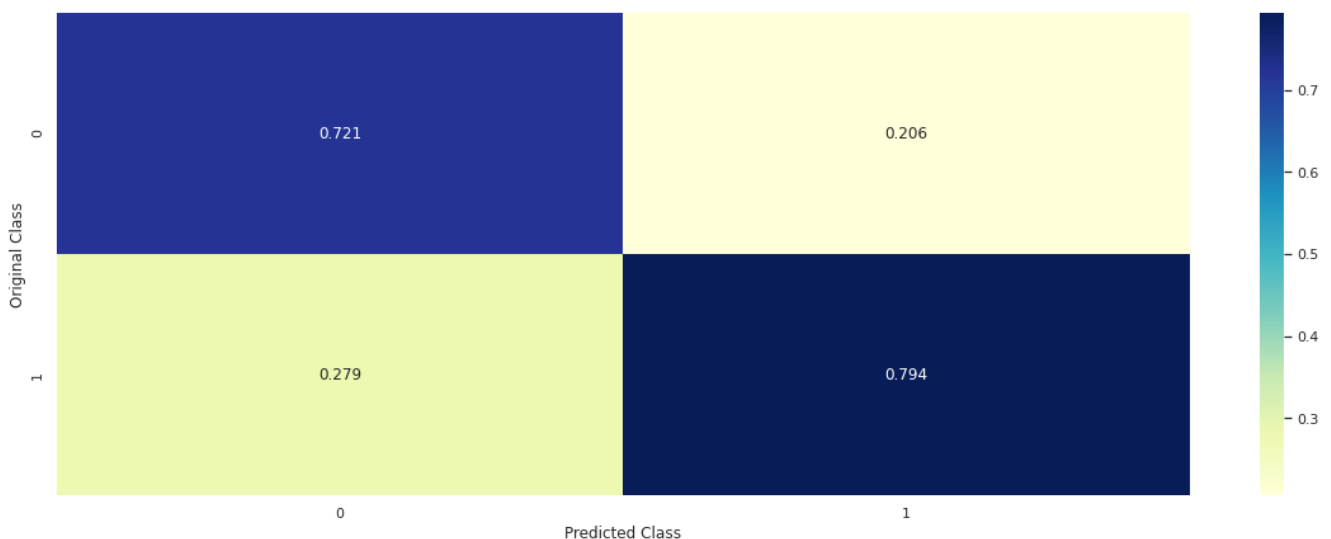
```
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
predict_and_plot_confusion_matrix(X_train, y_train, X_cv, y_cv, clf)
```

Log loss : 0.5217361177668973  
 Number of mis-classified points : 0.2478666027010918


----- Confusion matrix -----



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



- This was correctly classified before upsampling by all models

In [124]:

```
# from tabulate import tabulate
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(X_train,y_train)
test_point_index = 1
no_feature = 1000
predicted_cls = sig_clf.predict(Xr[test_point_index].reshape(1, -1))
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(Xr[test_point_index].reshape(1, -1)),4))
print("Actual Class :", yr[test_point_index].reshape(1, -1))
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,:no_feature]
```

```
Predicted Class : 0
Predicted Class Probabilities: [[0.8069 0.1931]]
Actual Class : [[1]]
```

**Lets summarize above models before proceeding with the feature engineering approach.**

In [125]:

```
from prettytable import PrettyTable

ptable = PrettyTable()
ptable.title = "*** Model Summary *** [Performance Metric: Log-Loss]"
ptable.field_names=["Model Name","Train LogLoss","CV LogLoss","Test LogLoss","% Misclassified Points"]
ptable.add_row(["Logistic Regression With Class balancing","0.298","0.297","0.302","0.122"])
ptable.add_row(["Linear SVM","0.309","0.308","0.312","0.122"])
ptable.add_row(["Random Forest Classifier ","0.265","0.266","0.272","0.095"])
ptable.add_row(["Logistic Regression With Class balancing(UPSAMPLING) ","0.522","0.521","0.522","0.247"])

print(ptable)
```

```
+-----+-----+-----+-----+
|               Model Name               | Train LogLoss | CV LogLoss | Test LogLoss |
| % Misclassified Points |               |             |              |
+-----+-----+-----+-----+
```

+-----+-----+-----+-----+							
-----+							
	Logistic Regression With Class balancing		0.298		0.297		0.302
	0.122						
	Linear SVM		0.309		0.308		0.312
0.122							
	Random Forest Classifier		0.265		0.266		0.272
0.095							
	Logistic Regression With Class balancing(UPSAMPLING)		0.522		0.521		0.522
	0.247						
+-----+-----+-----+-----+							
-----+							
◀					▶		

## CONCLUSION:

- All the models performed better than the random model, which makes sense.
- From the pretty table we can see that , **RandomForest** performed best here.
- Even the overfitting is not present if we check the train and test logloss, they are very close
- Over sampling method was also applied on the training data to make the data more balanced, but it gave worse results

In [ ]: